

Red Hat AI (RHAI) 3.5

New Key Features

Version 1, 2026-08-25

Home

Introduction

Course Title: Red Hat AI (RHAI) 3.5: New Key Features

Description: Welcome to the Red Hat AI (RHAI) 3.5 Enterprise Technical Training Course. This training guide covers the key architectural and functional updates in version 3.5.

This course is designed around the four core product pillars that drive the RHAI 3.5 strategy:

- Fast & Efficient Inference
- Connecting Data to Models
- Agentic AI
- Platform Infrastructure, Scale & Security

Each module represents a product pillar. Within each module, lessons are structured around the official Key Messages, with technical explanations, operational mechanics, and enterprise use-case scenarios.

Duration: >...

Technology Preview Features

Technology Preview features are not supported with Red Hat production service level agreements (SLAs) and might not be functionally complete. Red Hat does not recommend using them in production. These features provide early access to upcoming product features, enabling customers to test functionality and provide feedback during the development process. For more information about the support scope of Red Hat Technology Preview features, see [Technology Preview Features Support Scop](#).

Developer Preview Features

Developer Preview features are not supported by Red Hat in any way and are not functionally complete or production-ready. Do not use Developer Preview features for production or business-critical workloads. Developer Preview features provide early access to functionality in advance of possible inclusion in a Red Hat product offering. Customers can use these features to test functionality and provide feedback during the development process. Developer Preview features might not have any documentation, are subject to change or removal at any time, and have received limited testing. Red Hat might provide ways to submit feedback on Developer Preview features without an associated SLA. For more information about the support scope of Red Hat Developer Preview features, see [Developer Preview Support Scope](#).

Contributors

The PTL team acknowledges the valuable contributions of the following Red Hat associates:

- Sarvesh Pandit (Principal Technical Training Developer)
- Carlos Condado (Senior Product Marketing Manager, AI)
- Will McGrath (Senior Principal Product Marketing Manager, AI)
- Martin Isaksson (Principal Architect, AI)
- Younes Ben Brahim (Principal Product Marketing Manager, AI)
- Rutuja Deshmukh (Technical Training Developer)
- Anna Swicegood (Learning Product Manager)



Use of Artificial Intelligence: This publication contains content or media that is generated or assisted by AI tools. All the Training deliverables have primarily human authorship, and all content was reviewed by a human domain expert.

Fast & Efficient Inference

Production-ready agentic AI requires inference capabilities that go beyond simple text completion. Red Hat AI 3.5 gives you the operational depth, model safety transparency, resource governance, and multimodal foundations to deploy and manage high-priority LLM workloads at scale.

This section focuses on:

- Model Safety and Reliability Evidence Before Deployments
- The Operational Controls Shared Inference Requires
- The Foundation for AI Agents that Work Beyond Text

Model Safety and Reliability Evidence Before Deployments

The Operational Controls Shared Inference Requires

The Foundation for AI Agents that Work Beyond Text

Connecting Data to Models

Raw foundation models are functionally limited without access to enterprise data. RHAI 3.5 closes this gap with production-grade data integration pipelines, advanced model-tuning frameworks, and robust evaluation systems to build trustworthy AI applications.

This section focuses on:

- AutoRAG and AutoML: From Enterprise Data to Production AI
- Post-Training and Inference-Time Scaling: Improve Model Reasoning
- AI Evaluations and AI Safety: Build Confidence Before Production

AutoRAG and AutoML: From Enterprise Data to Production AI

Post-Training and Inference-Time Scaling: Improve Model Reasoning

AI Evaluations and AI Safety: Build Confidence Before Production

Practical Agentic AI

Red Hat AI 3.5 moves beyond basic model serving—delivering a governed, secure agentic platform built for production.

This section focuses on:

- Agentic API Foundation: OGX (GA)
- Pre-built Templates in AI Hub for Developer Productivity

Agentic API Foundation: OGX (GA)

Pre-built Templates in AI Hub for Developer Productivity

section3.adoc

section4.adoc

Enterprise Platform Infrastructure, Scale & Security

Scaling AI operations across your organization takes solid platform governance. Red Hat OpenShift AI 3.5 unifies safety, observability, multi-tenancy, and hardware acceleration into a single operational view.

This section focuses on:

- AI Security: Automated Red Teaming and Security Insights
- Measurable Enterprise AI: Zero-Configuration Observability
- Multi-Tenancy for Enterprise AI
- GPU-as-a-Service (GPUaaS)

AI Security: Automated Red Teaming and Security Insights

Measurable Enterprise AI: Zero-Configuration Observability

Appendix

FIXME: Content for Appendix...

Resources

[Download Course PDF](#)